

Introduction to Mobile Robotics

Bayes Filter – Particle Filter and Monte Carlo Localization

Wolfram Burgard, Michael Krawez

UTN

Particle Filter

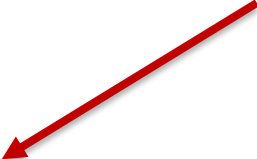
- Recall: Discrete filter
 - Discretize the continuous state space
 - High memory complexity
 - Fixed resolution (does not adapt to the belief)
- Particle filters are a way to **efficiently** represent **non-Gaussian distributions**
- Basic principle
 - Set of state hypotheses (“particles”)
 - Survival-of-the-fittest

Mathematical *Description*

- Set (actually a multi-set) of weighted samples



$$S = \left\{ \left\langle \underset{\substack{\uparrow \\ \text{state hypothesis}}}{s^{[i]}}, \underset{\substack{\uparrow \\ \text{importance weight}}}{w^{[i]}} \right\rangle \mid i = 1, \dots, N \right\}$$

- The samples represent the posterior $\begin{cases} 1 & \text{if } x=s^{[i]} \\ 0 & \text{otherwise} \end{cases}$

$$p(x) = \sum_{i=1}^N w_i \cdot \delta_{s^{[i]}}(x)$$


Function Approximation

- Particle sets can be used to approximate functions



- The more particles fall into an interval, the higher the probability of that interval

Bayes Filter with Particle Sets

- Measurement update

$$Bel(x) \leftarrow p(z|x) \overline{Bel}(x)$$

$$= p(z|x) \sum_i w_i \delta_{s^{[i]}}(x) = \sum_i p(z|s^{[i]}) w_i \delta_{s^{[i]}}(x)$$

- Motion update

$$\overline{Bel}(x) \leftarrow \int p(x | u, x') Bel(x') dx'$$

$$= \int p(x | u, x') \sum_i w_i \delta_{s^{[i]}}(x') dx' = \sum_i p(x | u, s^{[i]}) w_i$$

Particle Filter Algorithm

- Sample the next generation for particles using the proposal distribution
- Compute the importance weights :
$$weight = target\ distribution / proposal\ distribution$$
- Resampling: “Replace unlikely samples by more likely ones”

Particle Filter Algorithm

1. Algorithm **particle_filter**(S_{t-1}, u_t, z_t) returns S_t :
2. $S_t = \emptyset, \quad \eta = 0$
3. **For** $i = 1, \dots, n$ *Generate new samples*
4. Sample index $j(i)$ from the discrete distribution given by w_{t-1}
5. Sample x_t^i from $p(x_t | x_{t-1}, u_t)$ using $x_{t-1}^{j(i)}$ and u_t
6. $w_t^i = p(z_t | x_t^i)$ *Compute importance weight*
7. $h = h + w_t^i$ *Update normalization factor*
8. $S_t = S_t \cup \{x_t^i, w_t^i\}$ *Add to new particle set*
9. **For** $i = 1, \dots, n$
10. $w_t^i = w_t^i / h$ *Normalize weights*

Particle Filter Algorithm

$$Bel(x_t) = h p(z_t | x_t) \int p(x_t | x_{t-1}, u_t) Bel(x_{t-1}) dx_{t-1}$$

draw x_{t-1}^i from $Bel(x_{t-1})$

draw x_t^i from $p(x_t | x_{t-1}^i, u_t)$


Importance factor for x_t^i :

$$\begin{aligned} w_t^i &= \frac{\text{target distribution}}{\text{proposal distribution}} \\ &= \frac{h p(z_t | x_t) p(x_t | x_{t-1}^i, u_t) Bel(x_{t-1})}{p(x_t | x_{t-1}^i, u_t) Bel(x_{t-1})} \\ &\propto p(z_t | x_t) \end{aligned}$$


Resampling

- **Given**: Set S of weighted samples.
- **Wanted**: Random sample, where the probability of drawing x_i is given by w_i .
- Typically done n times with replacement to generate new sample set S' .

Resampling



- Roulette wheel
- Binary search, $O(n \log(n))$



- Stochastic universal sampling
- Systematic resampling
- Linear time complexity $O(n)$
- Easy to implement, low variance

Resampling Algorithm


1. Algorithm **systematic_resampling**(S, n):
2. $S' = \emptyset, c_1 = w^1$
3. **For** $i = 2 \dots n$ *Generate cdf*
4. $c_i = c_{i-1} + w^i$
5. $u_1 \sim U[0, n^{-1}]$, $i = 1$ *Initialize threshold*
6. **For** $j = 1 \dots n$ *Draw samples ...*
7. **While** ($u_j > c_i$) *Skip until next threshold reached*
8. $i = i + 1$
9. $S' = S' \cup \{ < x^i, n^{-1} > \}$ *Insert*
10. $u_{j+1} = u_j + n^{-1}$ *Increment threshold*
11. **Return** S'

Also called **stochastic universal sampling**

Mobile Robot Localization


- Each particle is a potential pose of the robot
- Proposal distribution is the motion model of the robot (prediction step)
- The observation model is used to compute the importance weight (correction step)

Motion Model Reminder




- Uncertainty in the translation of the robot:
Gaussian over the traveled distance
- Uncertainty in the rotation of the robot:
Gaussians over initial and final rotation
- For each particle, draw a new pose by sampling from these three individual normal distributions



Motion Model Reminder



Proximity Sensor Model Reminder




Laser



300cm

Sonar



400cm

Robot Localization using Particle Filters (1)

- Each particle is a potential pose of the robot
- The set of weighted particles approximates the posterior belief about the robot's pose (target distribution)

Robot Localization using Particle Filters (2)


- Particles are drawn from the motion model (proposal distribution)
- Particles are weighted according to the observation model (sensor model)
- Particles are resampled according to the particle weights

Robot Localization using Particle Filters (3)

Why is resampling needed?

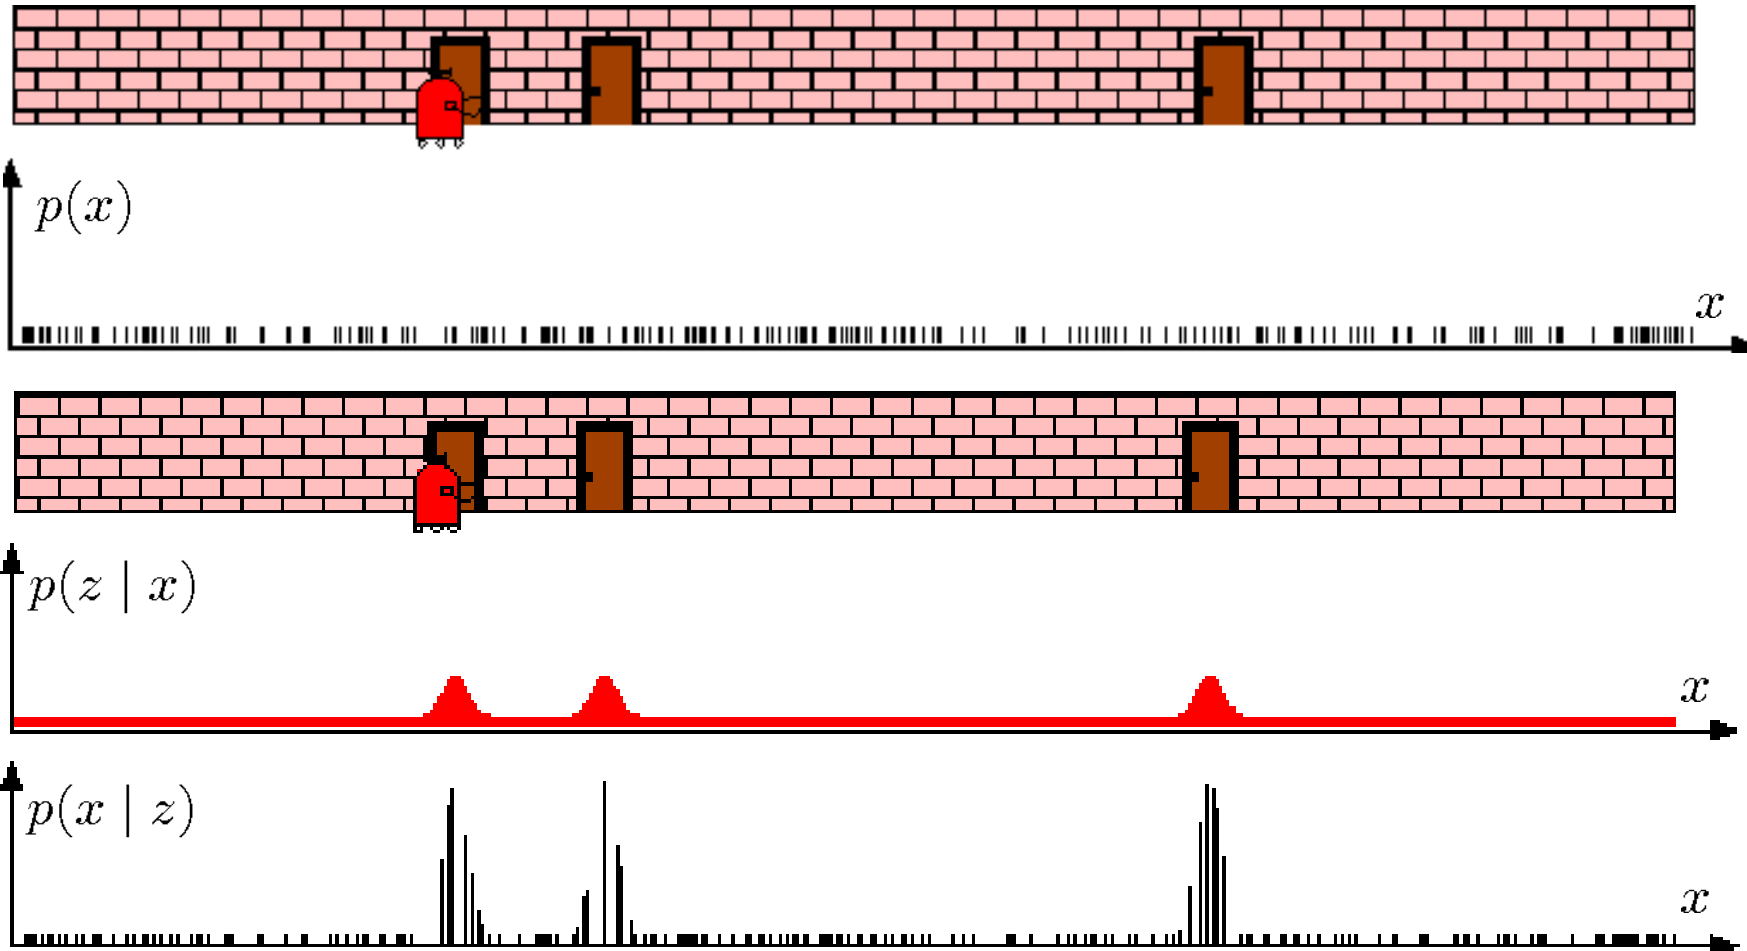
- We only have a finite number of particles
- Without resampling: The filter is likely to loose track of the “good” hypotheses
- Resampling ensures that particles stay in the meaningful area of the state space

Particle Filters





Sensor Information: Importance Sampling

$$\begin{aligned} Bel(x) &\leftarrow \alpha p(z | x) Bel^-(x) \\ w &\leftarrow \frac{\alpha p(z | x) Bel^-(x)}{Bel^-(x)} = \alpha p(z | x) \end{aligned}$$



Robot Motion

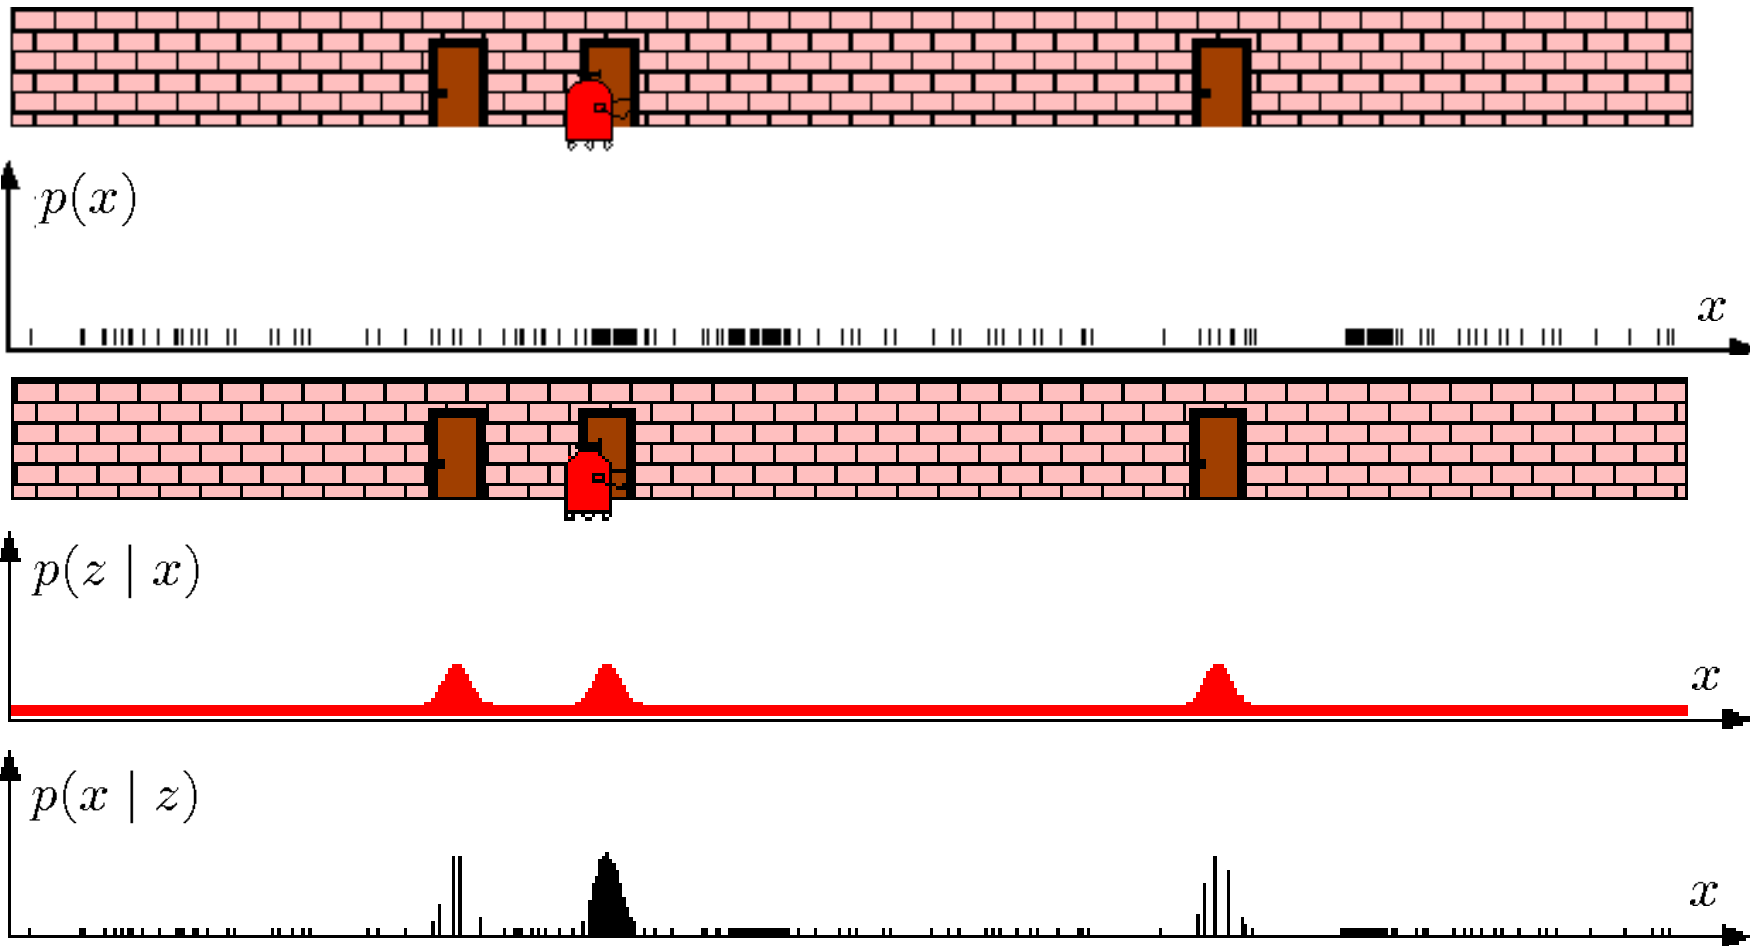
$$Bel^-(x) \leftarrow \int p(x | u, x') Bel(x') dx'$$



Sensor Information: Importance Sampling



$$Bel(x) \propto a p(z | x) Bel^-(x)$$

$$w \propto \frac{a p(z | x) Bel^-(x)}{Bel^-(x)} = a p(z | x)$$





Robot Motion


$$Bel^-(x) \propto \int p(x | u, x') Bel(x') dx'$$





Real World Application (LiDAR)



















































Highly Precise Localization and Positioning




Localization Accuracy





- KUKA omniMove (11t)
- Safety scanners
- Error in the area of millimeters
- Even in dynamic environments




Sample-based Localization (Ultrasound)




Discrete Filters Reminder (Ultrasound)



Initial Distribution




After Incorporating Ten Ultrasound Scans



After Incorporating 65 Ultrasound Scans








Estimated Path




Importance Sampling with Resampling: Landmark Detection Example





Distributions



Distributions






Wanted: samples distributed according to
 $p(x \mid z_1, z_2, z_3)$




This is Easy!


- We can draw samples from $p(x|z_i)$ by adding noise to the detection parameters.



Importance Sampling with Resampling




Weighted samples



After resampling

Localization for AIBO Robots



Limitations


- The approach described so far is able
 - to track the pose of a mobile robot and
 - to globally localize the robot
- How can we deal with localization errors (i.e., the kidnapped robot problem)?

Approaches

- Randomly insert a fixed number of samples with randomly chosen poses
- This corresponds to the assumption that the robot can be teleported at any point in time to an arbitrary location
- Alternatively, insert such samples inverse proportional to the average likelihood of the observations (the lower this likelihood the higher the probability that the current estimate is wrong).


Vision-based Localization (Recovery from Failure)

Random samples



Localization for AIBO Robots

Drawing from observations



Summary – Particle Filters

- Particle filters are an implementation of recursive Bayesian filtering
- They represent the posterior by a set of weighted samples
- They can model arbitrary and thus also non-Gaussian distributions
- Proposal to draw new samples
- Weights are computed to account for the difference between the proposal and the target
- Monte Carlo filter, Survival of the fittest, Condensation, Bootstrap filter

Summary – Particle Filter Localization

- In the context of localization, the particles are propagated according to the motion model.
- They are then weighted according to the likelihood model (likelihood of the observations).
- In a re-sampling step, new particles are drawn with a probability proportional to the likelihood of the observation.
- This leads to one of the most popular approaches to mobile robot localization